

The Top 20 Business Problems

That you can solve with Data Science.



Matt Dancho

Table of Contents

Table of Contents	2
Top 20 Business Problems (that can be solved with Data Science)	2
The Top 20 Business Problems Explained	6
Problem 1: Churn (Customer & Employees)	6
Problem 2: Time Series Forecasting	7
Problem 3: Lead Scoring (Email & Sales Marketing)	8
Problem 4: Production & Deployment	9
Problem 5: Customer Segmentation	10
Problem 6: Fraud Detection with Anomaly Identification	11
Problem 7: Market Basket Analysis (Recommendations)	12
Problem 8: Web-Scraping to Build a Competitor Database	13
Problem 9: A/B Testing (Website & Emails)	14
Problem 10: Marketing Channel Attribution (Budget)	15
Problem 11: Financial Investment & Portfolio Optimization	16
Problem 12: Social Sentiment Analysis for Brands	17
Problem 13: Network Analysis for Influential Customers	18
Problem 14: PDF Text Scraping and Document Analysis	19
Problem 15: Customer Lifetime Value (CLV & RFM)	20
Problem 16: Marketing Mix Modeling (MMM)	21
Problem 17: Automate Enterprise Data Pipelines with Big Data	22
Problem 18: Risk Analysis & Business Simulation	23
Problem 19: Probabilistic Business Prediction with Bayesian	24
Problem 20: Business Process Optimization (Linear & Nonlinear Constraints)	25
Problem 21: Geospatial Analysis (GIS)	26
Want to see how I solved the 75 most common business problems?	27

Top 20 Business Problems (that can be solved with Data Science)

Table 1 captures the **Top 20 Business Problems** that can be solved with data science and the R programming language. I will then discuss each more in depth with supporting information in Section 2.

Rank	Business Problem	Financial Impact	Algorithms Used	R Packages	Where to learn?
1	Churn (Customers & Employees)	\$\$\$	Classification (Logistic Regression)	tidymodels h2o	R-Track Program DS4B 1&2
2	Time Series Forecasting (Sales Prediction)	\$\$\$	Time Series Forecasting	modeltime timetk	R-Track Program Time Series
3	Lead Scoring (Email & Sales)	\$\$	Classification (Logistic Regression)	tidymodels h2o	R-Track Program DS4B 1&2
4	Production & Deployment	\$\$\$	Any	shiny plumber	R-Track Program Shiny 1&2
5	Customer Segmentation	\$\$	Clustering (K-Means) Dimensionality Reduction (PCA, UMAP)	stats umap	R-Track Program DS4B 1
6	Fraud Detection (Insurance Fraud, Credit Card Transactions)	\$\$\$	Anomaly Detection	H2O	Learning Labs PRO Lab 17
7	Intelligent Product Recommendations	\$\$	Association Rules	recommenderlab	Learning Labs PRO Lab 11
8	Build a Competitor Database	\$	Web Scraping	rvest	Learning Labs PRO Lab 8

Rank	Business Problem	Financial Impact	Algorithms Used	R Packages	Where to learn?
9	A/B Testing (Website & Email Marketing)	\$\$	Statistical Tests	tidyverse	Learning Labs PRO Lab 24
10	Marketing Channel Attribution	\$\$	Markov Models	ChannelAttribution	Learning Labs PRO Lab 25
11	Financial & Portfolio Analysis	\$\$	Optimization	tidyquant	Learning Labs PRO Lab 9
12	Sentiment Analysis	\$	Text Analysis & Natural Language Processing	tidytext	Learning Labs PRO Lab 32
13	Customer Network Analysis	\$\$	Network Analysis	tidygraph ggraph	Learning Labs PRO Lab 19
14	PDF Document Text Extraction	\$	Text Analysis & Image Processing	tidytext	Learning Labs PRO Lab 70
15	Customer Lifetime Value (CLV & RFM)	\$\$	Machine Learning Classification & Regression	tidymodels h2o	Learning Labs PRO Lab 58
16	Marketing Mix Modeling (MMM)	\$\$	Optimization	robyn	Learning Labs PRO Lab 61
17	Automate Enterprise Data Pipelines with Big Data	\$\$	Spark	sparklyr	Learning Labs PRO Lab 65
18	Risk Analysis & Business Simulation	\$\$	Simulation & Markov Chain Monte Carlo (MCMC) Sampling	tidyverse	Learning Labs PRO Lab 68
19	Probabilistic Business Prediction with Bayesian	\$\$	Bayesian	brms bayesian	Learning Labs PRO Lab 74 & 75
20	Business Process Optimization (Linear & Non-Linear Constraints)	\$\$	Optimization	ROI ompr	Learning Labs PRO Lab 15 & 16

Rank	Business Problem	Financial Impact	Algorithms Used	R Packages	Where to learn?
21	Geospatial Analysis (GIS)	\$\$\$	Networks	sf mapview	<u>Learning Labs PRO</u> Lab 76 & 77

The Top 20 Business Problems Explained

Problem 1: Churn (Customer & Employees)

Churn is when a customer or employee leaves the company. Churn is a very costly problem that eats away at the organization's revenue when customers leave or employees quit.

Churn is a **binary classification problem** where we are trying to determine whether or not something will happen. Hence either a 1 or a 0. However, the 1 or 0 is actually just a label that is based on a class probability (e.g. class 1 = 0.75 and class 0 = 0.25).



To solve this problem, I use:

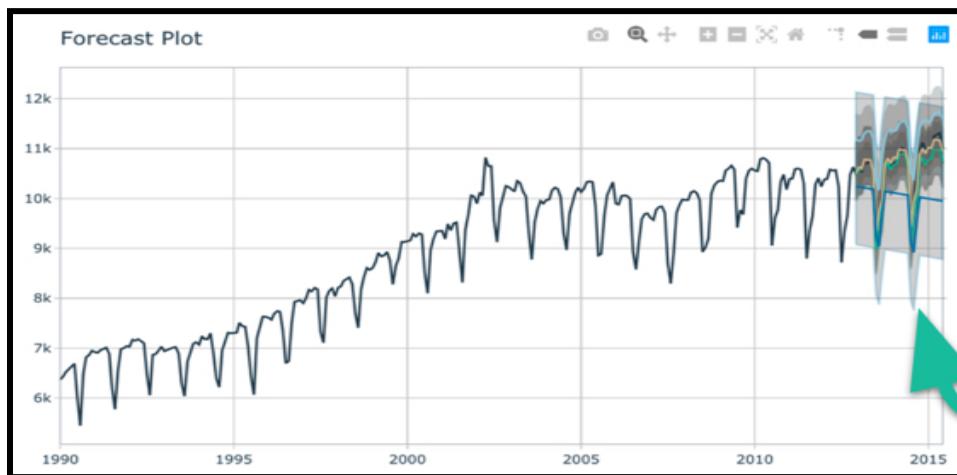
- **H2O Automatic Machine Learning** to predict the churn risk of each employee
- **Local Interpretable Model-Agnostic Explanations (LIME)** to explain why the model predicts the employee as likely to stay or leave the company.

I then make a **recommendation algorithm** to help organizations decide which actions to take. I cover how to solve this problem in-depth in my [5-Course R-Track Program](#).

Problem 2: Time Series Forecasting

Time series forecasting is the process of using historical data (such as sales revenue over time) as inputs to predict (estimate) the next H periods (often denoted the forecast horizon). Everything from hiring new employees and purchasing raw materials via the supply chain is driven off of a demand forecast. Therefore, organizations that improve their forecasting can save millions of dollars.

A demand forecast may look something like this for a *single time series*.



Companies need to make **hundreds of thousands of forecasts** every week or month (think of forecasting for every product your company makes). This is called **high-performance forecasting**.

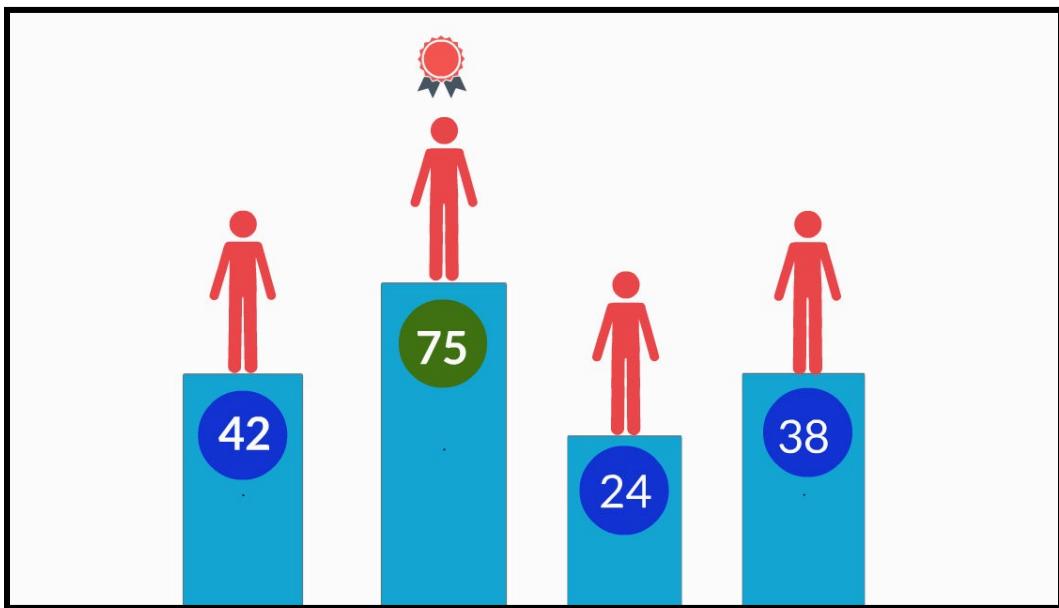
I cover high-performance forecasting using the winning strategies from 4 time series competitions using my **modeltime** & **timetk** R packages (forecasting R packages I created) in my [5-Course R-Track](#).



Problem 3: Lead Scoring (Email & Sales Marketing)

Lead scoring is a binary classification problem where you want to predict the probability of a lead turning into an order. Traditionally this is done by measuring the number of actions taken by a customer and then calculating a score by weighting each action in terms of the business's idea of how important the action is.

Data science can be used to more accurately predict the customers likelihood by combining the customer's actions taken (traditional method) and additional data like customer's time with the company, age, geographic region, and more. The result is a more accurate lead score.



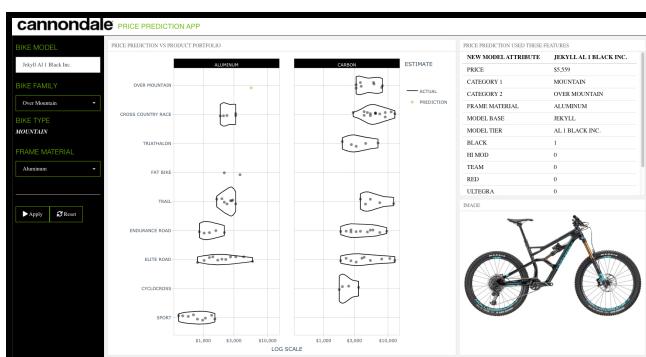
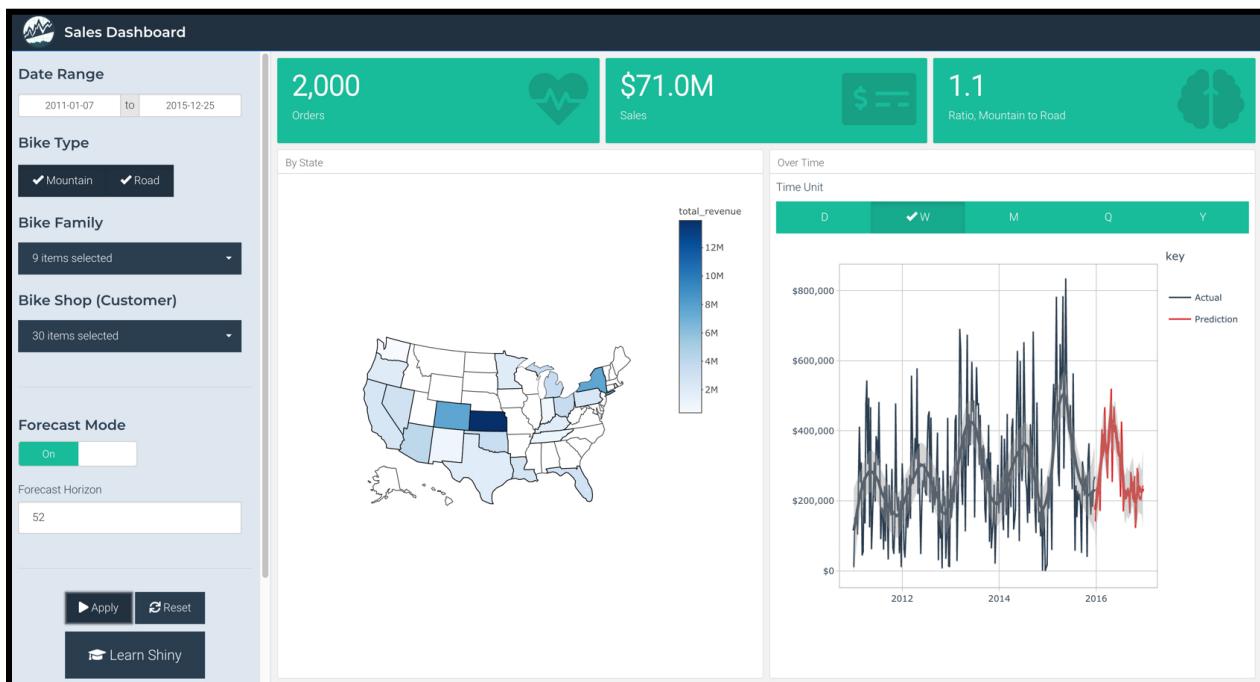
Lead scoring is a variant of Problem 1: Churn. The main difference is that instead of classifying the customer as Yes/No, we wish to **rank each of the customers** in order of most likely to purchase to least likely. Customers that are more ready to purchase can be marketed to, and those that are least likely should be nurtured.

I teach how to solve the binary classification problem you will face when you do lead scoring in my [5 Course R-Track](#). My preference is to use **H2O Automatic Machine Learning and LIME for explainable machine learning**.

Problem 4: Production & Deployment

Production is the process of taking a machine learning algorithm or decision-making analysis code (script) and converting it into a form that is usable by the business. For R, the most common methods are to build and deploy shiny apps or plumber APIs.

The most common production app is a dashboard. The advantage of using Shiny is that a predictive analysis that is codified in R can quickly be converted into an app. Think of the app as the package for the code, and users interact with the app (using dropdowns and buttons), which then automates the process of running the code behind the scenes. Shiny apps are covered in **Shiny Web Apps Parts 1 & 2 Courses**.

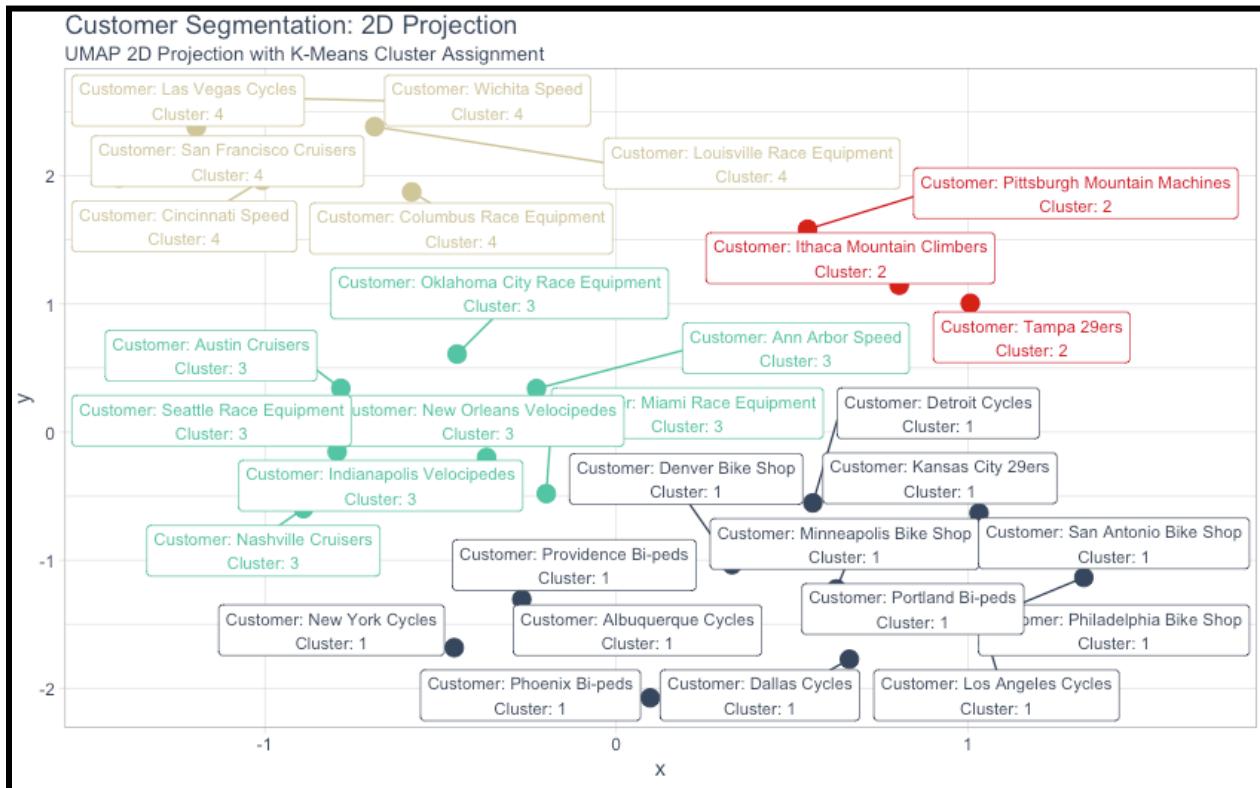


In my [5-Course R-Track](#), you learn how to build 4 shiny web applications (2 shown here) that cover the most common scenarios:

- Pulling data from APIs,
- integrating predictive models, and
- Displaying visualizations (images, maps, and statistical / time series plots) like those shown.

Problem 5: Customer Segmentation

Customer segmentation is the process of separating customers into groups based on common purchasing behavior (e.g. this group tends to buy more shoes vs this group tends to buy more handbags). Companies can then target these segments and tailor marketing messages most commonly via email. **Segmentation increases revenue and improves customer satisfaction. Win-win.**



In my 5-Course R-Track, you use clustering and dimensionality reduction to visualize segments of customers. I show you how to apply:

- **K-Means:** A common clustering algorithm that does well at identifying segments within business data
- **UMAP:** A dimensionality reduction technique that allows us to plot the variance of the data in two dimensions

The result is a customer segmentation that can be used to target similar customers with tailored marketing messages.

The customer segmentation problem is solved in my [5-Course R-Track program](#).

Problem 6: Fraud Detection with Anomaly Identification

Fraud detection is identifying abnormal transactions or behavior from customers. Often the abnormal behavior can indicate that something isn't right about the transaction or insurance claim, and therefore requires a hold until further evaluation can take place.

Detecting Fraudulent Transactions

Key Issues

- What does **abnormal behavior** mean?
- How do we handle **Big Data Sets**?
- Which techniques work gracefully in the presence of **High Imbalance**?

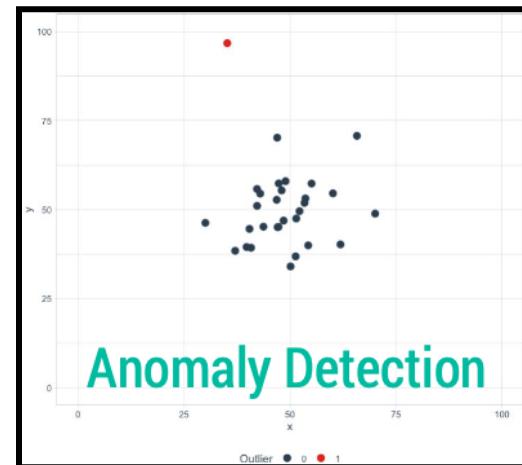
```
> # 1.1 CLASS IMBALANCE ----  
> credit_card_tbl %>%  
+   count(Class) %>%  
+   mutate(prop = n / sum(n))  
# A tibble: 2 x 3  
  Class     n    prop  
  <dbl> <int> <dbl>  
1     0 284315 0.998  
2     1    492 0.00173
```

A common method is **anomaly detection** which is used to identify abnormal events. This can be interpreted as suspicious behavior.

I use **H2O Isolation Forests** for anomaly detection using the h2o package in R.

For **time series anomaly detection**, I created the anomalize R package. I have since incorporated most of the time series anomaly detection functionality into my **timetk R package**.

Learn how I solved this problem in Learning Lab 17, available in [Learning Labs PRO](#).



Problem 7: Market Basket Analysis (Recommendations)

Market Basket Analysis is a technique used to identify which products get purchased together. We can then use this information to recommend products based on a customer's cart (prior to checkout). This tactic increases the Average Order Value (AOV).

Generating **Business Value** with Market Basket Analysis



Business Objectives

- Understand what customers buy
- Improve purchasing experience
- Better configure product placement
- Target customers for upsell

Market Basket Analysis
Can help



One of the most common techniques I use to mine customer cart's and make recommendations is **association rules**. I use the arules R package to implement various association rules mining.

Learn how I solved this problem in [Learning Lab 11](#).

Problem 8: Web-Scraping to Build a Competitor Database

Building and tracking your competitors can help companies better understand competitor behavior and look for areas of improvement to seek as a competitive advantage in the marketplace.

The internet has a massive amount of product and service information, and **many of your company's competitors list this information online**. Your organization can take advantage of the free data by building a strategic database. The question is then, how best to collect the competitor data? That's where web-scraping comes in.

Why Web Scrape?

Product Pages have a ton of Valuable Information

COMPONENT	DETAILS
FORK	Lefty LightPip w/ integrated SuperHesus LED Light
BRAKES	Shimano M316 Hydro Disc
CRANK	FSA Comet, Gates Belt Drive
WHEELSET	Bad Boy 27.5 Rims / Lefty Hubs
COMPONENTS	Bad Boy Stem/Fabric Silicone Grip, Crank Coop Radius Saddle/integrated LED Seatpost
EXTRAS	N/A
FRAMES	Bad Boy OPI, SmartForm C1 Alloy, integrated Urban Armor bumper, BSA, 1-1/8" headtube, sliding dropout
FORK	Lefty LightPip rigid, integrated SuperHesus LED lighting, USB rechargeable battery, 1-1/8" steerer
RIMS	Bad Boy C1, 650b, double wall, 36 spokes
HUBS	Lefty 10 front, Shimano Alfine 8-speed internal rear
SPOKES	Stainless Steel, 14g

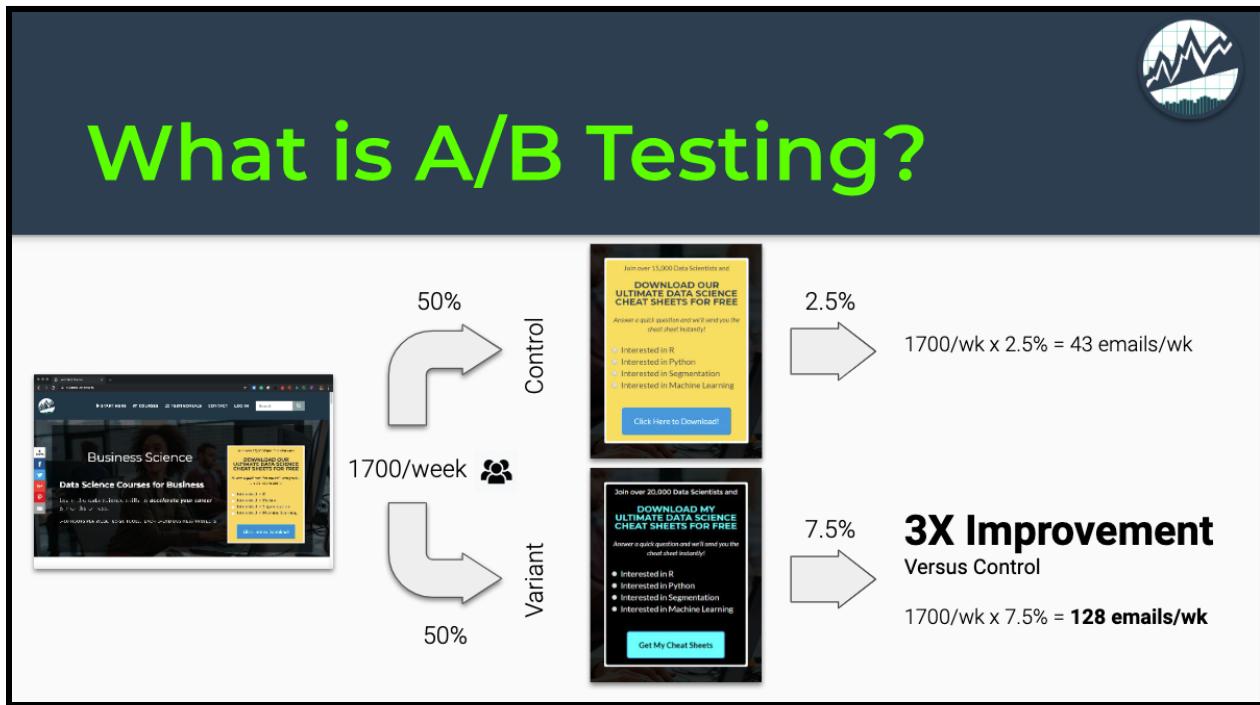
Web-scraping is the most common method for developing a database of competitive information. In R, I commonly use the:

- rvest R package to extract HTML data into R
- tidyverse to convert the raw HTML into tabular data

Learn how I solved this problem in [Learning Lab 8](#).

Problem 9: A/B Testing (Website & Emails)

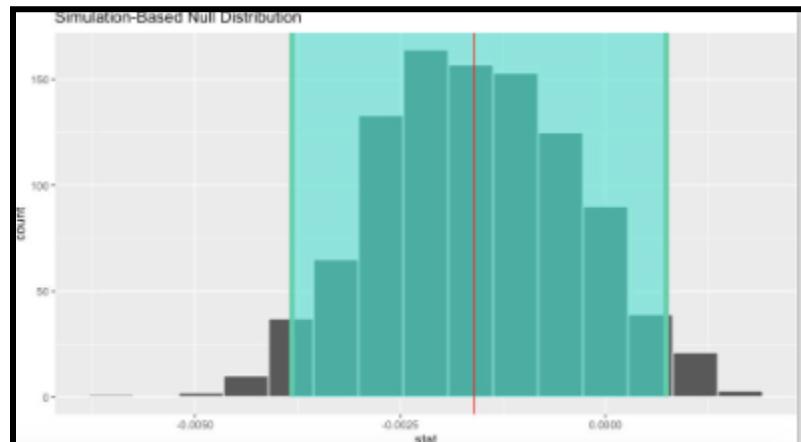
A/B testing is the process of testing changes to websites and emails by dividing the traffic into two and randomly assigning to two variants. It's most common to make small changes to each variant, and then to test the impact using statistical analysis.



In R, I use the **tidyverse** to analyze web-traffic and email split tests.

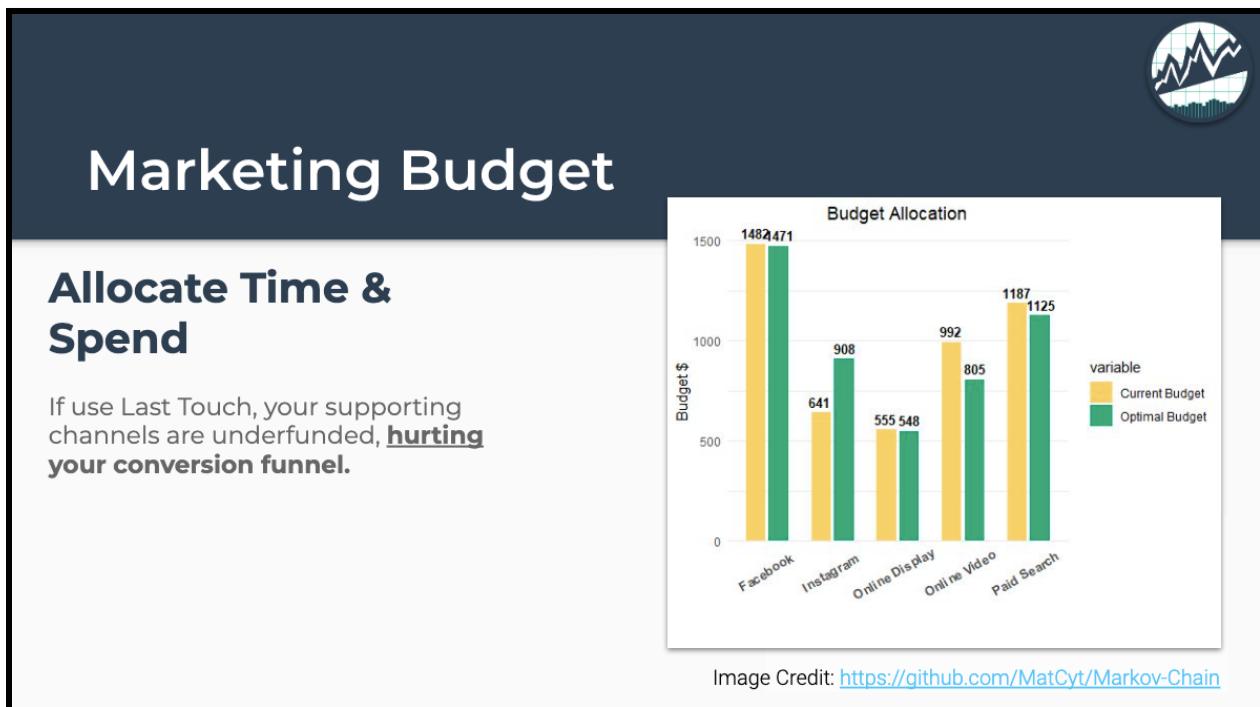
Then I add **confidence intervals** to my assessments to determine if we should accept or reject the null hypothesis.

Learn how I solved this problem in [Learning Lab 24](#).



Problem 10: Marketing Channel Attribution (Budget)

Marketing Attribution (Channel Attribution) is a common problem that marketers face when allocating the correct amount of money to each marketing advertisement channel (e.g. Facebook, Instagram, Paid Search) to optimize for conversions.



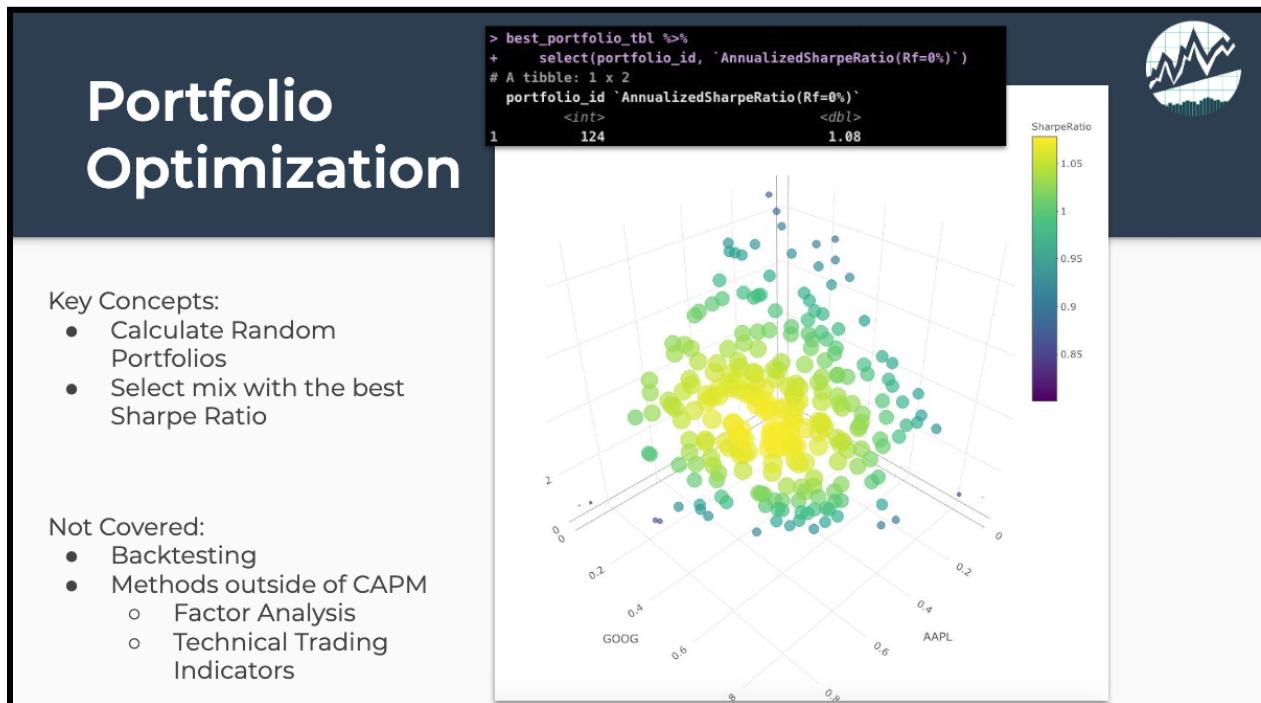
I use the **ChannelAttribution** R package in R to help solve the Marketing Attribution Problem.

A newer technique is the **FB Robyn algorithm (robyn package)** is available for Marketing Mix Modeling (MMM) (See Problem 16), a similar and related business problem.

Learn how I solved this problem in [Learning Lab 25](#).

Problem 11: Financial Investment & Portfolio Optimization

Financial investment and portfolio analysis is the process of aggregating stock data, using technical indicators, making portfolios and optimizing to reduce key metrics like the Sharpe Ratio in order to improve investment decision-making.



I created the **tidyquant R package**, which provides a streamlined workflow for financial investment and portfolio analysis.

Tidyquant has:

- An API to pull financial data into R
- A full suite of financial and technical analysis tools
- Integrated portfolio analysis for combining financial assets and analyzing a portfolio

Learn how I solved this problem in [Learning Lab 9](#).

Problem 12: Social Sentiment Analysis for Brands

Sentiment Analysis is used to assess the positive or negative score of text. One key application is in the use of **Brand Sentiment**, which can be used for companies like Coca-Cola to analyze social media sentiment as a gauge for how well their brand is being viewed to the public. Positive sentiment indicates your customers are viewing your brand positively, and negative sentiment could be cause for concern.

Positive Sentiment

Row ID: 36
Screen Name: CropAndSoil
Text:
Congratulations to Michelle Schroeder-Moreno for winning the NC State faculty #Sustainability Award for her work developing our agroecology program! Recipients showed achievement in sustainability on campus, the community or globally. Awards sponsored by @CocaCola. Great job!
<https://t.co/EApHEK3Ppg>

Coca-Cola

Sustainability

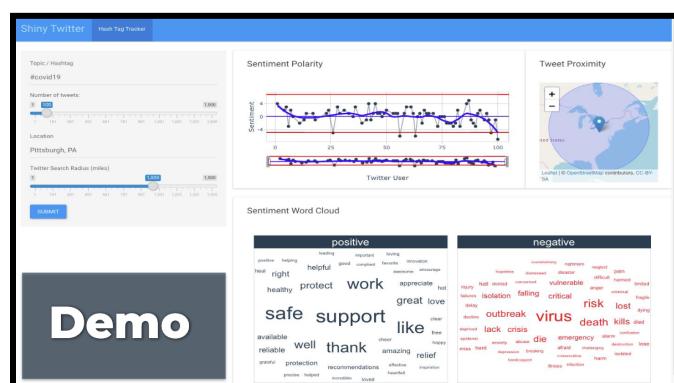
Community

Positive Sentiment
Can highlight where people are positively viewing your brand. Keep this up!

This shiny app I made combines:

- Rtweets Twitter API
- Tidytext Text Analysis
- Tidyverse Data wrangling & Viz
- Shiny Web Apps

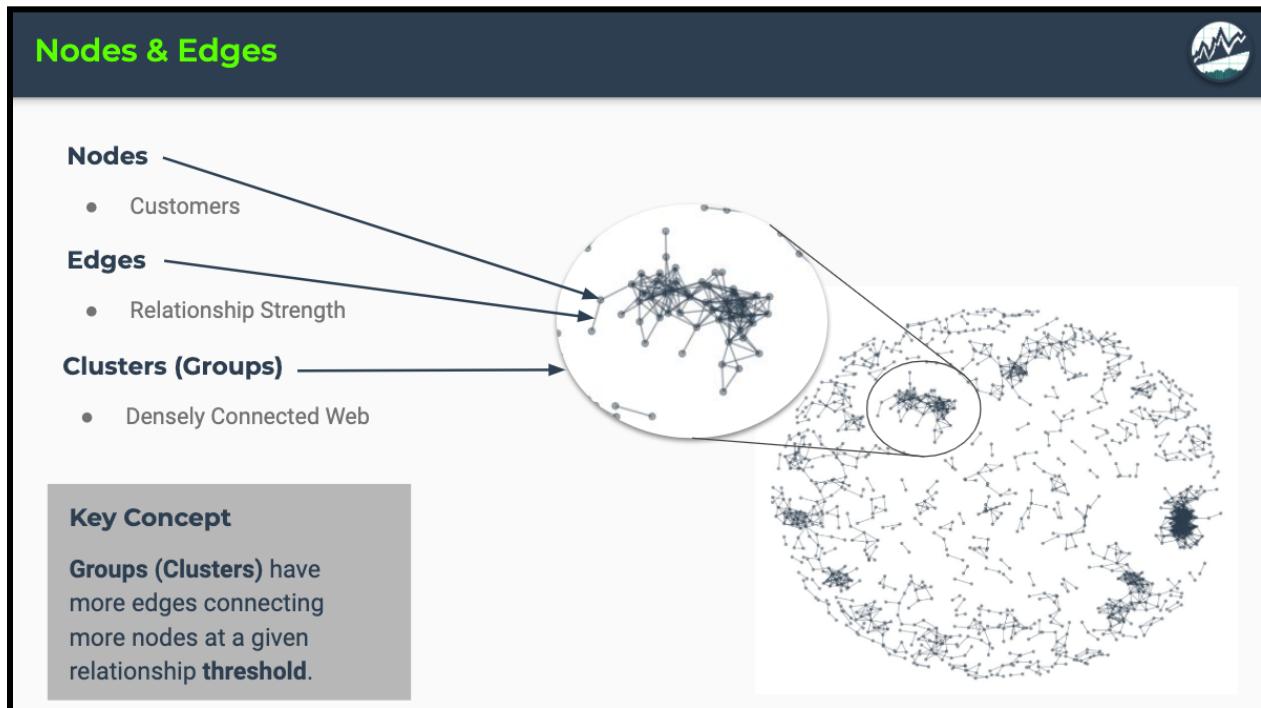
To make a Shiny App for Sentiment analysis for Brand Assessment.



Learn how I solved this problem in [Learning Lab 32](#).

Problem 13: Network Analysis for Influential Customers

Network analysis can be used to identify influential customers. The key difference between network analysis and clustering is that in clustering you just have a group assignment label, but in network analysis you gain the ability to understand which customers are related AND which customers are the most influential. Once you know which customers are most influential, you can design marketing messages to the “influencers”, which will have a higher effect on the entire group.



I use these two R packages for network analysis:

- **tidygraph** for network analysis
- **ggraph** for network visualization

Both of these packages follow the **tidyverse-centric concepts** learned in my 5-Course R-Track. Once you learn the tidyverse, these network analysis R packages are easy to pick up.

Learn how I solved this problem in [Learning Lab 19](#).

Problem 14: PDF Text Scraping and Document Analysis

PDF and Document Text-Scraping is becoming increasingly popular. Businesses have mounds of unstructured data stored as PDF, Word, and Excel documents. This data can be mined for important information.

One example is **Resume Analysis**. I use a technique called Natural Language Processing (NLP) to evaluate the similarity of a job applicant's resume to the company's job description. This allows me to filter out unqualified candidates, without wasting time reviewing thousands of resumes.

The screenshot shows a shiny application interface. At the top, there is a logo of a line graph. Below it, the title "NLP: Text Analysis of Resumes" is displayed. On the left side of the interface, there are three main sections: "What skills does the candidate have?", "Resume Analysis", "What skills does the job require? Job Description Analysis", and "Is the candidate a good fit? Similarity". On the right side, there is a profile section for "Matt Dancho" (Data Scientist, Trainer & Mentor), a "Work Experience" section listing his roles at Bonney Forge and Data Science, and a "Term Frequency" cloud visualization.

What skills does the candidate have?
Resume Analysis

What skills does the job require? Job Description Analysis

Is the candidate a good fit? Similarity

Matt Dancho
Data Scientist, Trainer & Mentor

Work Experience

Founder
Business Science
February 2010 - Present

- Identified gap in the market for data science courses that teach applied skills to solve real world business problems that led to multimillion dollar business
- Scalped business to 3000+ students in first 3 years
- Increased revenue 20% Year-over-Year
- Used NLP and predictive lead scoring to segment email list leading to 30% subscriber growth
- Sample Project 1: Pricing Recommendation Application (R, Shiny)
- Sample Project 2: Time Series Forecasting Application (R, Shiny)

Director, Product Engineering & Market Development
Bonney Forge
February 2017 - 2018

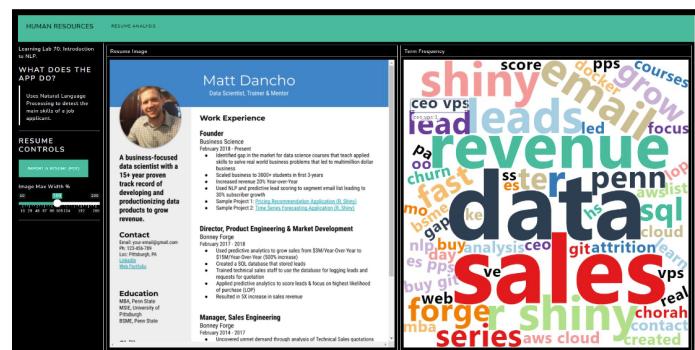
- Used predictive analytics to grow sales from \$3M/Year Over Year to \$15M/Year-Over-Year (50% increase)
- Created a SQL database that stored leads
- Trained technical sales staff to use the database for logging leads and requests for quotation
- Resulted in 30% increase in sales revenue

Term Frequency

This shiny app I made combines:

- shiny for web applications
- tidytext for text analysis
- magick for images
- tidyverse for similarity

To make a **resume analyzer** web application.



Learn how I solved this problem in [Learning Lab 70](#).

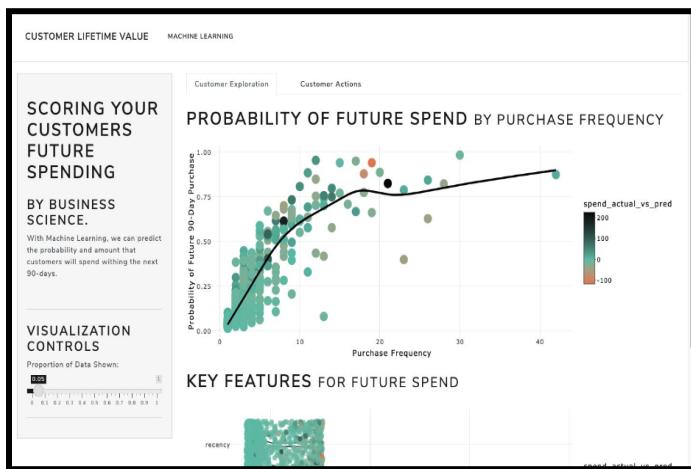
Problem 15: Customer Lifetime Value (CLV & RFM)

Customer Lifetime Value (CLV) is used to highlight which customers to focus on. Companies use CLV to estimate the profitability of the future relationship with a customer. Most algorithms use **RFM (Recency-Frequency-Monetary)** attributes for customers to help estimate customer lifetime value.

There are two approaches to CLV including (1) traditional economic and (2) machine learning approaches. I favor machine learning as I normally get higher accuracy and better predictive insights when I model using the newer technology (machine learning).

Customer Lifetime Value (CLV)

Companies use this as a metric to gauge profitability and to **focus on which customers**. In short, CLV is the **profit** from estimated by the future relationship with a customer. There are **many different approaches** to modeling CLV.



This shiny CLV app I made combines:

- **shiny** for Web Apps
- **tidymodels** for CLV estimation using Machine Learning

The customer lifetime value app **scores the customers** by estimated future purchasing in the next 90-days.

Learn how I solved this problem in [Learning Lab 58](#).

Problem 16: Marketing Mix Modeling (MMM)

Marketing Mix Modeling (MMM) is a technique that helps marketers quantify the impact of advertising and marketing actions on the customer revenue that's generated. MMM is often used in the **marketing budget allocation process** to account for causation in the key marketing variables such as paid media variables, organic variables, contextual variables, and time-based and seasonal variables.

The infographic features a large purple funnel centered on a white background. At the top right of the slide is a circular icon containing a line graph with a peak and a bar chart below it. The main title "Goal: Increase ROI" is at the top left. Below the title, a text box contains the following information:

Organizations have a conversion funnel that generates **ROI from advertising**.

Hypothetical Example:

4 / 1 = 400% ROI

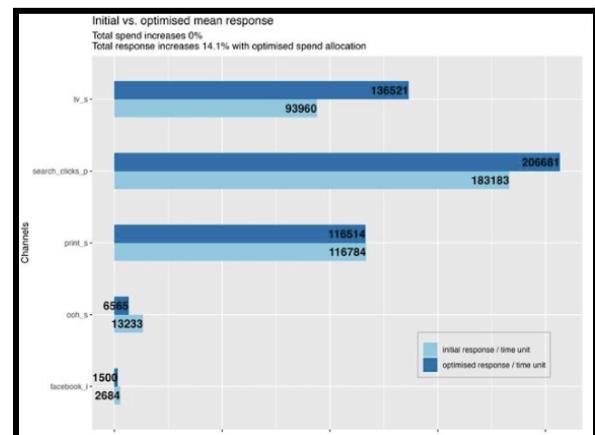
To the right of the funnel, there are two points connected by a line:

- A purple dot labeled "\$1 In" with a line pointing to the top of the funnel.
- A purple dot labeled "\$4 Out" with a line pointing to the bottom of the funnel.

Below the funnel, the text "Customer Revenue" is aligned with the bottom point.

Here I'm using the **Robyn R package** developed by Facebook to help *automate* the MMM process, thus reducing the human bias in budget allocation decisions.

Learn how I solved this problem in [Learning Lab 61](#).



Problem 17: Automate Enterprise Data Pipelines with Big Data

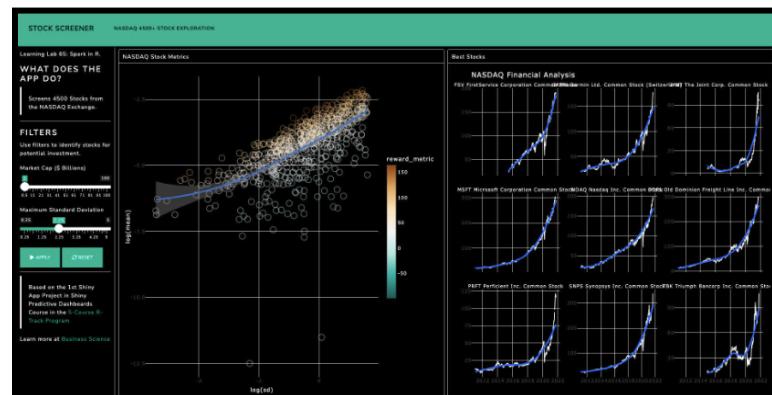
Data is getting bigger to the point where it becomes difficult for people to process enterprise data on their computers. To meet these big data demands, organizations are adopting tools like Spark, which is a unified analytics engine for large-scale data processing.

The slide has a black header bar with the title "What is Spark?" in green. In the top right corner of the header is a small circular icon containing a line chart. The main content area contains a quote in white text: "A unified analytics engine for large-scale data processing." Below the quote is the Apache Spark logo, which consists of the word "APACHE" in a smaller serif font above the word "Spark" in a large, bold, sans-serif font. A stylized orange star shape is positioned to the right of the "Spark" text.

This shiny app I made combines:

- **shiny** for Web Apps
- **Spark** for Big Data (**sparklyr**)

To automate the processing of 4,500 stocks I'm evaluating on the NASDAQ stock exchange.

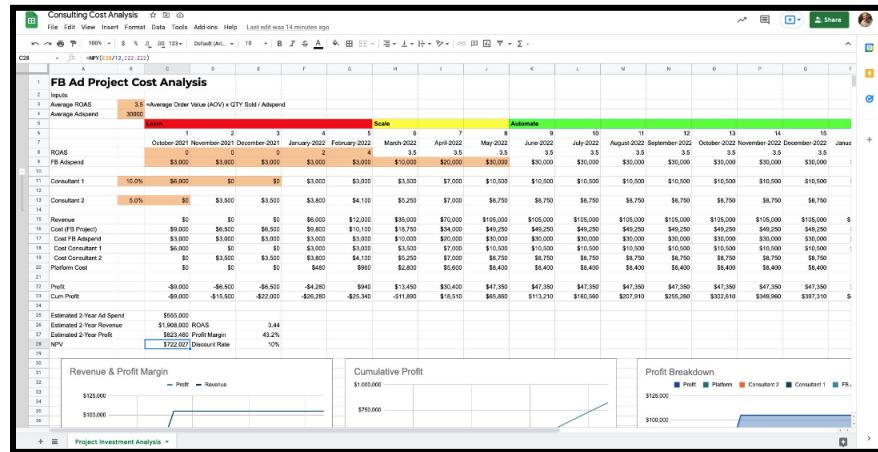


Learn how I solved this problem in [Learning Lab 65](#).

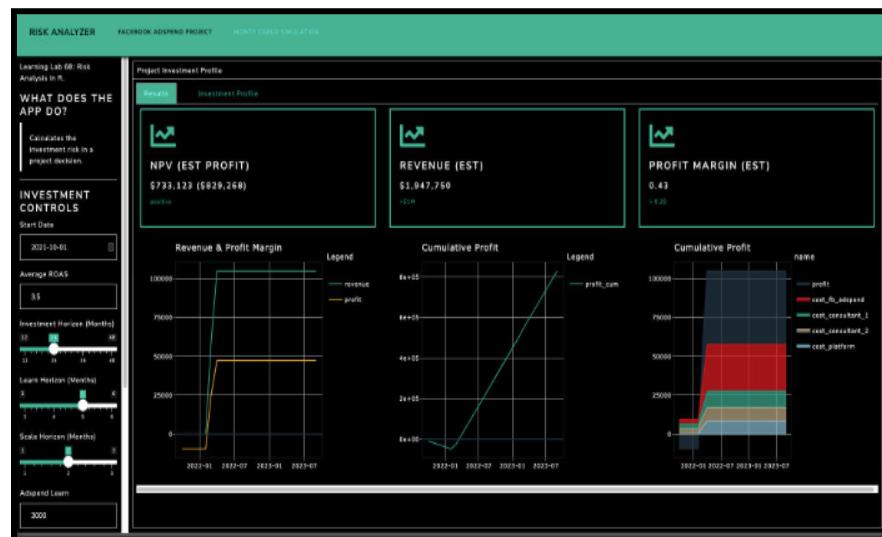
Problem 18: Risk Analysis & Business Simulation

Business analysis and simulation is done in Excel using expensive add-ons like the Palisade Decision Tools suite. But, R is a free replacement, which can save hundreds of thousands of dollars in expensive software, reducing errors, and automating the risk analysis process.

R can be used to convert this Excel Capital Expenditure Investment Project from a single Net Present Value (NPV) Calculation into a full **Risk Simulation with confidence intervals**.



Next, we can build a shiny app to empower the business and automate business simulations.



I solved this problem in [Learning Lab 68](#).

Here I'm using:

- The **tidyverse** for simulation
- **shiny** for web applications

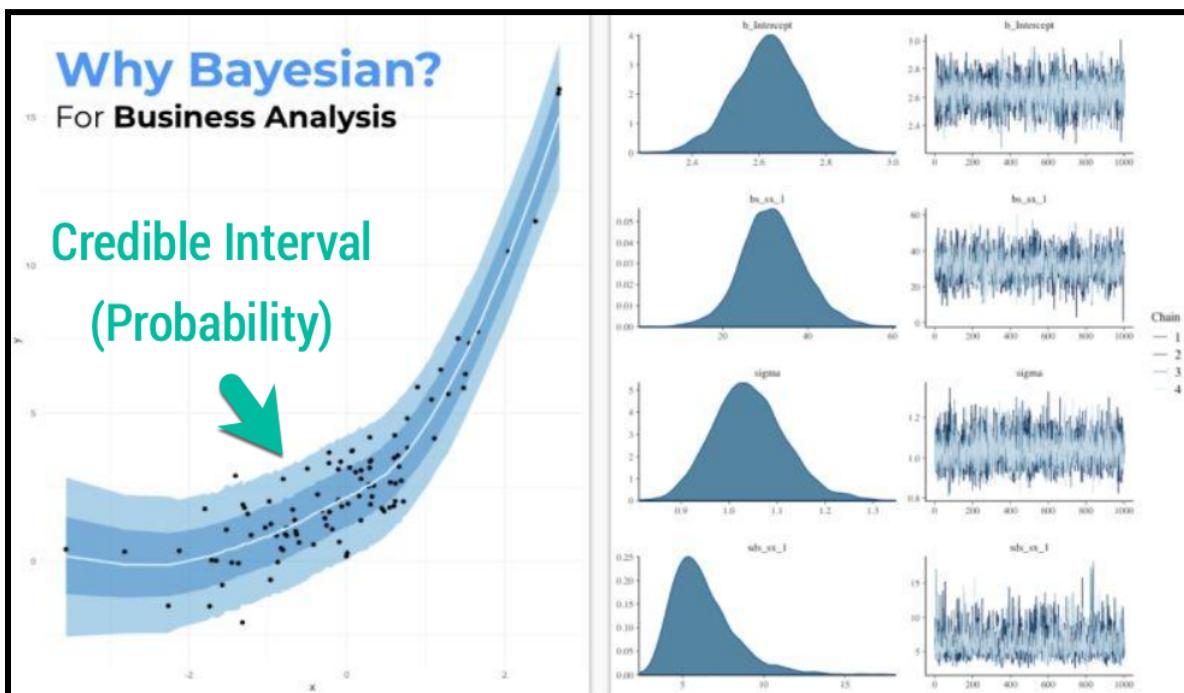
The result is a shiny web application that users can use to **automate business simulation and risk analysis**.

Problem 19: Probabilistic Business Prediction with Bayesian

In prediction, many times it's more important to **understand your confidence** than simply make a prediction. Understanding confidence allows organizations to make better decisions that reduce their *downside risk*, a key factor in determining the profitability of a model.

This is where **Bayesian methods** can help us to predict with a more accurate probability estimate. Bayesian uses a unique approach called Markov Chain Monte Carlo (MCMC) resampling to estimate a Credible Interval.

The Bayesian Credible Interval is a probability that surrounds the prediction. This differs from the traditional approach by adjusting to the variance of the data rather than assuming a constant variance (which is incorrect in most cases).



For bayesian analysis and probabilistic prediction, I use the **brms R library** and a **tidymodels extension named bayesian**.

I solved this problem in [Learning Labs 74 & 75](#).

Problem 20: Business Process Optimization (Linear & Nonlinear Constraints)

Business process optimization is used to optimize an objective (e.g. maximize profit) given that a company has limited resources (e.g. production line constraints, labor constraints, supply constraints).



Business Process Model

Labor Costs	Cost per labor hour assembling \$14 Cost per labor hour testing, line 1 \$22 Cost per labor hour testing, line 2 \$19
Process Attributes	Inputs for assembling and testing a Macbook
Manufacturing Plan	Model 1 Model 2 Model 3 Model 4 Model 5 Model 6 Model 7 Model 8
Sales Estimates	Labor hours for assembly 4 5 5 5 5.5 5.5 5.5 6 Labor hours for testing, line 1 1.5 2 2 2.5 2.5 2.5 2.5 3 Labor hours for testing, line 2 2 2.5 2.5 2.5 3 3 3.5 3.5 Cost of component parts \$900 \$1,350 \$1,350 \$1,350 \$1,500 \$1,500 \$1,500 \$1,800 Selling price \$2,100 \$2,700 \$2,700 \$2,820 \$3,000 \$3,150 \$3,180 \$3,600 Unit margin, tested on line 1 \$1,111 \$1,236 \$1,296 \$1,356 \$1,368 \$1,518 \$1,548 \$1,650 Unit margin, tested on line 2 \$1,106 \$1,233 \$1,293 \$1,353 \$1,366 \$1,516 \$1,537 \$1,650
Labor Constraints	Assembling, testing plan (# of Macbooks)
Objective	Model 1 Model 2 Model 3 Model 4 Model 5 Model 6 Model 7 Model 8 Number tested on line 1 1500 0 0 0 0 100 1000 0 Number tested on line 2 0 0 0 0 0 900 0 500 Total computers produced 1500 0 0 0 0 1000 1000 500 Maximum sales 1500 1250 1250 1250 1000 1000 1000 800 Constraints (hours per month) Hours used Hours available Labor availability for assembling 20000 <= 20000 Labor availability for testing, line 1 5000 <= 5000 Labor availability for testing, line 2 4450 <= 6000 Net profit (\$ per month) Model 1 Model 2 Model 3 Model 4 Model 5 Model 6 Model 7 Model 8 Totals Tested on line 1 \$1,666,500 \$0 \$0 \$0 \$0 \$0 \$151,800 \$1,548,000 \$0 \$3,366,300 Tested on line 2 \$0 \$0 \$0 \$0 \$0 \$0 \$1,364,400 \$0 \$824,750 \$2,189,150 \$5,555,450

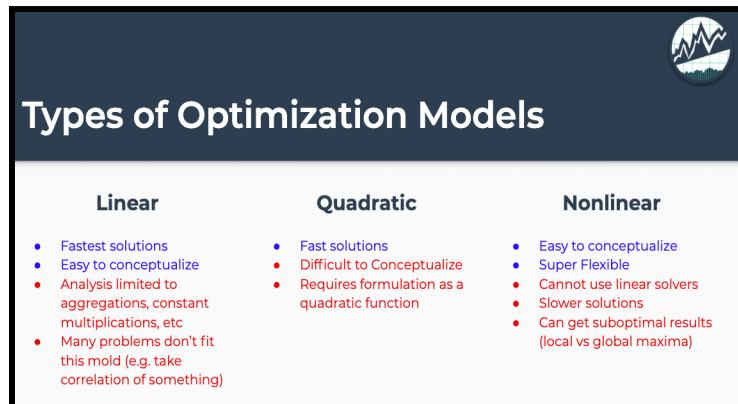
The two R packages that I use are:

- **ompr:** linear constraints ($y = mx$)
- **ROI:** non-linear constraints ($y = mx^2$)

Implementing these two methodologies, I've optimized problems in R including:

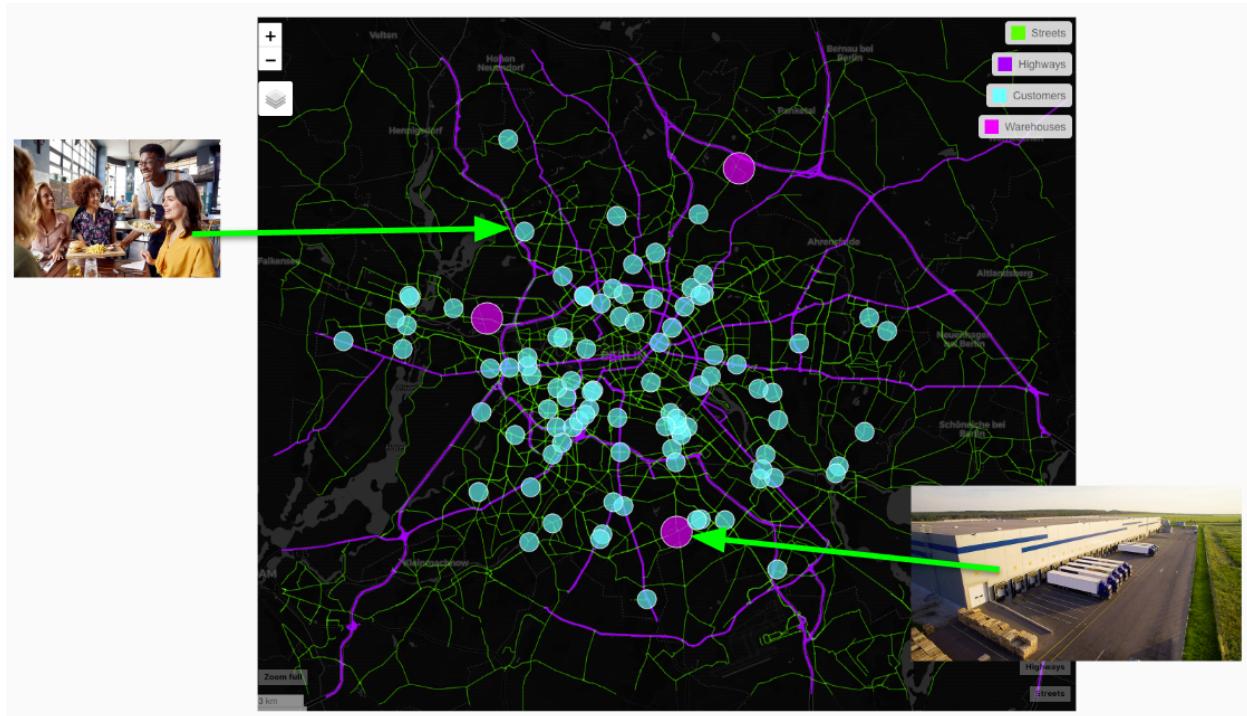
- **Linear:** Business processes to determine the optimal number of products to maximize profit
- **Nonlinear:** The optimal stock portfolio mix to maximize the Sharpe ratio

I solved this problem in [Learning Labs 15 & 16](#).



Problem 21: Geospatial Analysis (GIS)

Geospatial analysis is becoming increasingly important to businesses that seek to use geographic information to help better plan and connect with customers. One area of increasing importance is to reduce transportation costs and better plan distribution locations based on customer density.



The main R packages that I use are `sf` and `mapview`. The `sf` package provides the data structure and data wrangling for geospatial data including geospatial joins and other important operations. The `mapview` package makes it easy to produce geospatial visualizations.

If you'd like to learn more, I've solved the transportation routing problem and the customer-distributor problems in [Learning Labs 76 & 77](#).

Want to see how I solved the 75 most common business problems?

[Join Learning Labs PRO](#)

Lab 34: Advanced Customer Segmentation & Market Basket App (E-Commerce)

Advanced Customer Segmentation Python & R

Matt Dancho & David Curry
Business Science Learning Lab

Learn from Matt's unique portfolio of over 75 projects that companies need.

Geospatial (1+ Labs)
Bayesian (2 Labs)
NLP & Text Analysis in R & Python (3 labs)
Excel Simulation in R & Python (2 labs)
Spark in R & Python (3 labs)

Time Series in R & Python (2 labs)
Time Series Forecasting (8+ labs)
Customer Analytics: R Shiny | Python Dash (2 Labs)
R Targets in Production (2 labs)
NEW Tidymodels Extensions (6 labs)
Building an R Package (3 labs)
R in Production (4 labs)
Reticulate Python + R Series (5 Labs)
Shiny + APIs (5 labs)
Marketing Analytics (4 Labs)
SQL for Data Science (3 Labs)

Over 75+ Labs!

[See what Labs are available inside Learning Labs PRO](#)